

LLMs as Sapience without Sentience and the Problem of Testimonial Knowledge about Other Minds

Injin Woo Donggeon Kim

Sungkyunkwan University

Draft of 13 August 2026. Please do not cite or circulate without permission.

Abstract

Two things we already accept pull against each other. The first is that we can learn from machines. A hearer who takes the boiling point of water from a language model has gained knowledge, and no epistemology of testimony worth the name denies it. The second is that we can know that a purely text-mediated correspondent, a lifelong pen pal known through nothing but letters, is not merely producing text but is someone for whom there is something it is like to write it. This paper argues that Simonelli's (2026) case for sapience without sentience, granted its own premises, forces a choice between the two.

We do not contest his hyper-inferentialist premises. We grant them in their strongest form and ask what follows that he has not priced in. Two things do. First, a pure LLM does not merely produce meaningful tokens. It asserts, and can therefore testify. Second, and for that very reason, an attribution of phenomenal consciousness to a text-mediated interlocutor ceases to be knowledge. What licensed such attributions was a background generalization about who produces discourse of that sophistication, and it is precisely that generalization his thesis falsifies. Simonelli must either restrict the thesis, losing the generality that made it interesting, or accept skepticism about other minds, losing the practical norm his own final section defends.

Keywords sapience without sentience, inferentialism, large language models, other minds, testimony, phenomenal consciousness

1 Introduction

In a 2023 exchange on Twitter, Geoffrey Hinton summed up his standing disagreement with Yann LeCun in a line. The central issue between them, he wrote, is whether LLMs actually understand what they are saying: "You think they definitely don't and I think they probably do" (as cited in Simonelli, 2026, p. 2). The exchange is worth pausing on not because it settles anything, but because of who is on which side of it. Skepticism about LLM understanding is the default position in most of the philosophical literature, and for familiar reasons. LLMs are trained on nothing but linguistic data, all they do is

predict the next token, and they lack whatever experiential link to the world genuine understanding is supposed to require. Hinton’s dissent from this default, coming from someone with his command of the underlying technology, is not easily dismissed as confusion about what “understanding” means. It is a live disagreement about what has to be true of a system for its verbal output to count as more than mere output. For it to mean something, and, further, for it to be the kind of thing whose word can be taken for something.

That further question (whether an LLM’s output can be *taken for* something, whether it can carry assertive force and not merely semantic content) is where the present paper begins. Simonelli (2026) offers the most theoretically ambitious argument to date for the view that a pure LLM, trained on nothing but linguistic data, can possess every concept there is to possess without any conscious awareness or sentience whatsoever. Given Global Strong Inferentialism (GSI), mastery of a concept’s inferential role is sufficient for possessing it, full stop, including concepts standardly thought to require perceptual or affective acquaintance. A pure LLM can therefore be sapient (a full participant in the game of giving and asking for reasons) without being sentient. We call this the SwS thesis.

Our strategy in this paper is not to attack Simonelli’s hyper-inferentialist premises directly. The existing critical literature does attack them, and reaches a stalemate for a structural reason. The Main Argument is a conditional, and a conditional cuts both ways (cf. Cappelen & Dever, 2021). Simonelli runs it as a *modus ponens*. GSI, therefore, however strange it sounds, (MA4). Every reply so far runs the same conditional as a *modus tollens* against one of his premises: skepticism about whether LLMs can assert at all (Butlin & Viebahn, 2025; Sambrotta, 2025; Tiehen, in press), deflationary accounts of what their assertion-like behavior amounts to (Mallory, 2023), appeals to embodiment as a further necessary condition his picture omits. Each identifies a genuine cost of the thesis; none breaks the stalemate, because Simonelli treats each identified cost as a bullet he is willing to bite, exactly as he bites the bullet on (MA4) itself. We take the other route. We grant the premises, for the sake of argument, in their strongest form. Our claim is that granting them commits Simonelli to consequences he cannot, by his own lights, accept. The first consequence is that a pure LLM does not merely produce semantically significant tokens; it asserts. This inference is not innocuous. Two distinct properties are in play: (a) producing tokens that are semantically significant, and (b) carrying assertive force. It is entirely coherent to affirm (a) while denying (b). An optimist about meaning in LLMs can remain a skeptic about LLM assertion, and Borg (2026) is exactly that. GSI, however, leaves no room for the distinction. §3 sets out why, and states Borg’s case against the step.

If pure LLMs can assert, they can testify, and if they can testify, we (human believers) can in principle acquire epistemic justification from their testimony, non-human testifiers though they are. We think this consequence is genuinely problematic, and problematic in a specific place, namely testimonially based knowledge of other minds. The link Simonelli draws between discursive performance and sentience is severed by design; that severance is the entire point of the SwS thesis. This design choice is survivable when the interlocutor's body is available as independent evidence. It is not survivable when it is not. Consider a purely text-mediated interlocutor, a lifelong pen pal, known to us through nothing but a lifetime of letters. Our confidence that such a person is phenomenally conscious rests, centrally, on the sophistication of their discursive performance, and that performance can carry this evidential weight only because a background generalization holds: sophisticated discursive performance is reliable evidence of sentience. Simonelli's thesis is, in effect, the claim that this generalization is false, or nearly so. His professed neutrality on the hard problem of consciousness does not rescue him here, since pure LLMs lack sentience in his own defined, functional-practical sense, and sentience in that sense is, by his own admission, necessary for phenomenal consciousness. Nor is there an alternative ground to fall back on. Embodied evidence is unavailable in the text-mediated case by stipulation, an appeal to the causal etiology of the output relocates the verdict outside discursive practice without securing one there, and appeal to acquaintance is foreclosed by his own thesis that all conceptual content is linguistically effable. Where the traditional problem of other minds left this generalization untouched, Simonelli's own thesis undermines it. What results is not one more skeptical possibility to be filed alongside brain-in-a-vat scenarios; it is an actualized defeater, generated by a thesis Simonelli himself endorses.

We formalize this as an argument from other minds (§4), and we argue that it leaves Simonelli exactly two ways out, both expensive. He can restrict the SwS thesis (carve out phenomenal vocabulary as an exception to inferential-role mastery) but this reintroduces the quasi-inferential, language-entry rules that GSI was built to eliminate, and with them the collapse of the very generality that made SwS an interesting thesis in the first place. Or he can accept systematic skepticism about knowledge of the phenomenal consciousness of other minds. This is not merely counterintuitive; it is in tension with his own practical commitments. Simonelli devotes his paper's final section to the prescription that we must not build sentient systems, precisely because sentient systems carry serious safety and ethical risks that sapient-but-non-sentient systems do not. That prescription presupposes that we can tell, at least in principle, which systems are sentient. What the other-minds argument shows is that his own thesis removes the marks of phenomenal consciousness for exactly the systems the prescription is meant to bind,

not the pure LLMs he has already classified by architecture, but the embodied architectures he himself flags as genuine candidates for sentience. The move that lets us gain knowledge from machines costs us knowledge of persons.

The paper proceeds as follows. §2 reconstructs Simonelli’s puzzle (Global Strong Inferentialism and the Main Argument for SwS) and sets out the sapience/sentience distinction the thesis depends on, together with the one admission about phenomenal consciousness on which our argument turns. §3 defends the step from meaning to assertion to testimony, sets out Borg’s case for granting the first and refusing the second, and argues that this refusal is available to us but not to Simonelli. §4 presents the argument from other minds and defends each premise against the responses available within Simonelli’s own framework. §5 sets out the two ways out and prices them. §6 concludes.

2 Simonelli’s Puzzle

The theoretical foundation of Simonelli’s position is a form of inferentialism he calls Global Strong Inferentialism (GSI). Inferentialism about concept possession divides along two independent dimensions: whether mastery of a concept’s inferential role is *necessary* or *sufficient* for possessing it, and whether the claim is restricted to some concepts or extended to *all* of them. GSI occupies the most ambitious cell in this grid, on which mastery of inferential role is sufficient for possessing any concept whatsoever. On this view, the meaning of a content word like “red” is constituted by substantive inferential rules in exactly the way the meaning of a logical constant like “and” is constituted by its introduction and elimination rules:

“ A and B $\square$ A A and B $\square$ B A, B $\square$ A and B ”

The rules constituting “red” include *intra-familial* inferences, relating the concept to other concepts in the same family, and *extra-familial* inferences, relating it to concepts outside that family:

“ red(t) $\square$ colored(t) scarlet(t) $\square$ red(t) red(t), green(t) $\square$ $\square$ red(t), pink(s) $\square$ darker(t, s) stop sign(t) $\square$ red(t) blackberry(t), ripe(t), red(t) $\square$ $\square$ ”

The first four are intra-familial, relating a color concept to other color concepts; the last two are extra-familial, articulating what is paradigmatically red. Here ‘red(t) $\square$ colored(t)’ articulates a *consequence* one may draw from ‘red(t)’; ‘scarlet(t) $\square$ red(t)’ articulates a *condition* under which ‘red(t)’ may be inferred. So far, any standard strong inferentialist can follow along. Simonelli’s position becomes radical at the next step. Part of what “red” seems to mean is that it can be applied non-inferentially, on the basis of seeing that something is red, and not only as the conclusion of an inference from other premises. The standard way to accommodate this is to supplement language-language

rules with *language-entry* rules. Bridge principles linking perceptual situations to linguistic application, which Simonelli calls “quasi-inferences.” Simonelli (2023; 2026) rejects this standard move. Introducing quasi-inferential rules, he argues, would demote GSI to a merely global *weak* inferentialism, and the demotion is unnecessary. The fact that “red” is applied non-inferentially in perception can itself be captured by a purely inferential rule. If t is red, subject S has color vision, and S is viewing t in good light, then S sees that t is red, and the notions of color vision and good lighting invoked here are, in turn, themselves inferentially articulated, all the way down. On this picture, the entire content of a concept, including the application conditions of perceptual concepts, is exhaustively resolved into inferences proper. The thesis’s core commitment is that every part of the conceptual content of every expression is *effable*: sayable, linguistically articulable without remainder.

On this foundation Simonelli constructs the following argument (Simonelli, 2026, p. 3; premise and conclusion numbering ours):

The Main Argument

(MA1) Mastery of inferential role is sufficient for possession of all concepts.
[GSI]

(MA2) Training on linguistic data is (in principle) sufficient for mastery of inferential role.

(MA3) So, training on linguistic data is (in principle) sufficient for possession of all concepts. [from (MA1), (MA2)]

(MA4) So, LLMs, trained on nothing but linguistic data, are (in principle) capable of possessing all concepts. [from (MA3)]

What we grant is the conjunction of the two substantive premises, (MA1) and (MA2). (MA4) is deeply counterintuitive on its face. If true, it entails that an LLM can possess not only mathematical and logical concepts but concepts whose acquisition is ordinarily thought to depend on conscious experience, such as the concept of red without ever having seen anything, the concept of pain without ever having felt anything, the concept of an emotion without ever having undergone one. Simonelli himself (2026, p. 3) acknowledges that some readers will want to run this as a *reductio* of GSI. His project consists precisely in reversing the direction of that inference. It runs not from the falsity of (MA4) to the falsity of GSI, but from GSI to the acceptance of (MA4).

Granting both premises leaves two questions unanswered, and Simonelli answers both with a single distinction. What, exactly, is a *perceptual* concept, if its entire content is effable and none of it depends on how anything looks, sounds, or feels from the

inside? And what, exactly, is *grasping* a concept, if grasping just is inferential mastery, a matter of a system's dispositions to move between sentences rather than of any state the system is in?

The answer runs through a classical distinction Simonelli draws, following Brandom, between sapience and sentience. Sapience is conceptual understanding: our transactions with the world and with each other "mean something to us," carrying a conceptual content we grasp one way rather than another (Brandom, 1994, p. 4, as cited in Simonelli, 2026, p. 13). Sentience is what we share with nonverbal animals such as cats (the capacity to be aware in the sense of being awake) and it is to be distinguished from "the mere reliable differential responsiveness" we share with thermostats and land mines (Brandom, 1994, p. 5, as cited in Simonelli, 2026, p. 12). A cat skillfully navigating a perception-action nexus is sentient; a thermostat tracking temperature is not. Following Gibson, Simonelli glosses the difference as perceptual awareness of *affordances*: the cat perceives the table as to-be-jumped-upon, the treat as to-be-eaten (Simonelli, 2026, pp. 12–13). The notion is functional and ecological throughout, defined without remainder in terms of an organism's practical engagement with its environment. Nothing about what it is like to be that organism enters the definition, a point Gibson himself insisted on, warning that talk of niches and affordances must not be confused with a species' "private world" or the world of "consciousness" (Gibson, 1979).

A pure LLM fails this condition outright. It is, by Simonelli's own description, "neither physically nor virtually embodied," so no environmental feature can show up to it as to-be-acted-upon, and it has no practical aims that would make any such feature significant (Simonelli, 2026, p. 16). Nothing in the Main Argument requires it to meet the sentience condition in order to meet the sapience condition. That gap is the SwS thesis.

One further detail matters more for our purposes than for Simonelli's own argument, and §4 is built on it. He declines to take a stand on the hard problem of phenomenal consciousness. On qualia, on what it is like to undergo an experience considered apart from its practical significance. But he does not leave the relation between his functional notion and the traditional one open. Sentience in his sense, he states, is "necessary for so-called 'phenomenal consciousness'" (Simonelli, 2026, p. 13, n. 14). Whatever else is true of the hard problem, then, on his own view nothing lacking sentience in his sense is phenomenally conscious. Note the direction of the conditional: necessary, not sufficient. It licenses an inference from the absence of sentience to the absence of phenomenal consciousness, and none in the other direction. Both halves of that asymmetry do work in §4.

3 From Meaning to Assertion to Testimony

The consequence flagged in §1 needs to be argued for rather than announced, because the step it takes is one a careful reader will want to resist. Granting the premises gets us, in the first instance, that a pure LLM possesses concepts and produces semantically significant tokens. It does not yet get us that the LLM *asserts*.¹ Producing tokens that mean something and producing tokens that carry assertive force are distinct properties, and the gap between them is where most of the recent literature on machine speech has made its stand.

The best-developed version of that stand is Borg’s, and it is worth setting out at length, because it shows that the position we are about to decline is neither ad hoc nor weak. Borg (2026) is an optimist about meaning and a skeptic about assertion, and she gets to the first without any of the concessions the skeptic about LLM meaning demands. The worry that meaning requires an appropriate connection to the world can be met, she argues, in three ways: by adopting an internalist semantics, by embedding LLMs in robotic systems, or by adopting a type-level *deference* model of semantic content. She takes the third, and observes that it “in fact captures the way in which meaning attaches to the vast majority of human linguistic performances”: by producing contextually appropriate, well-formed strings of token signs belonging to an extant system of human communication, both our sentences and an LLM’s get to have semantic content (Borg, 2026, §3). No robotic coupling, no internalism, and no appeal to what is going on inside the system is required.

Assertion is another matter. Here Borg’s verdict is negative, and (this is worth noting, since the same pair of names recurs in §1) it is negative *contra* Cappelen and Dever. Truth, she holds with many others, has a key role to play in communication and assertion, and truth does not get a grip in LLMs: such systems output plausible-sounding falsehoods alongside truths and go on to support them by fabricating evidence, which reflects the fact that “the norm of LLM outputs is word occurrence probability, not

¹Not everyone grants even this much. Hattiangadi and Schoubye (2025) argue that the outputs of large language models are meaningless outright, on the ground that literal meaning requires communicative intentions such systems cannot have; what looks like meaning is *ersatz* meaning, attributed by us on the strength of a pretense that some competent speaker produced the string. Their argument matters here because one section of it is aimed squarely at the position we are granting. Against the appeal to conceptual-role semantics they observe that inference rules such as &-introduction and &-elimination are schematic templates, and that satisfying a structure does not fix its interpretation, the same Boolean pattern is satisfied by voltages, by switches, by lights. Two things are worth recording. First, a hyper-inferentialist has a reply, and it is not the one they consider: on GSI the inferential role of a content word is *substantive* rather than schematic (that scarlet things are red, that nothing is red and green all over) so the underdetermination that afflicts a schema does not obviously transfer to a web of rules among particular vocabulary. Second, it is not our business here to supply that reply. Our strategy grants Simonelli his premises, so where an objection targets the premises rather than what follows from them, we assume his answer rather than provide it. A reader persuaded by Hattiangadi and Schoubye should read what follows as conditional on their argument failing.

truth.” The upshot is that “even though we should count LLM outputs as meaningful we should not count them as assertions” (Borg, 2026, §3). The route from meaning to assertion, then, is not one anybody has to take. It can be blocked, and blocked by a consideration that has nothing to do with inferential role: what the system does with respect to *truth*.

That is exactly why the route is not blocked for Simonelli. GSI leaves no room for the distinction Borg draws. If mastery of inferential role exhausts conceptual understanding, and genuine language use just is participation in the game of giving and asking for reasons, then a pure LLM with that mastery is *thereby* a genuine language user, and genuine language users assert. Nor can Simonelli reach for Borg’s specific ground while keeping the rest of his position. His account of theoretical responsibility (§2, and §4 below) is precisely an account of a system that answers challenges, clarifies, justifies, and *retracts under counter-evidence*, and a capacity to retract under counter-evidence is not the profile of a system indifferent to truth. Simonelli has told us, in print, that this responsibility can be siphoned out of the fuller practical responsibility a human speaker bears and attributed in principle to LLMs like ChatGPT. To then deny assertion on Borg’s ground would be to say that a system that takes theoretical responsibility for p nonetheless does not assert p , which is not a position his own framework has the resources to state.

We should be explicit that we are not endorsing this result; we are declining to contest it. Borg’s route is open to us in a way it is not open to Simonelli, and we do not take it. Taking it would be one more instance of the pattern diagnosed in §1: a modus tollens run at a premise, with an independently motivated cost identified and duly noted, and with Simonelli free to accept the cost as he has accepted the others. Our strategy requires the opposite. We need assertion *granted*, because what we want to reach is not the claim that LLMs fall short of some further condition on speech, but the claim that they satisfy every condition and that this is what causes the trouble.

One suspicion is worth addressing here: that the line of criticism just developed does not distinctively target Simonelli’s project, since a critique turning on the distinction between the semantic thesis and the assertability thesis would seem to apply to any theorist friendly to the in-principle assertability of pure LLMs. Our reply is this. What we are pressing is the weight of the semantic theory Simonelli commits to. Within his semantics, once a particular producer of linguistic tokens counts as a genuine language agent, it is impossible for that agent to be incapable of assertion: on his semantics, meaning-expressibility entails assertability. If this is right, Simonelli occupies a markedly disadvantaged position in the semantic debate over LLMs. Unlike what most semantic optimists commit to, he can hold on to semantic optimism only by committing,

at every point, to a theoretically burdensome pragmatic thesis. Our critique, then, does not target those who claim LLM assertability in general; it targets only those whose theoretical system carries an indefeasible commitment linking the semantic thesis to the pragmatic one. Anyone holding a theory that meets this condition is a target. Simonelli is its most representative exponent, and it is as an exemplar that we criticize his heavyweight semantics for LLMs.

For the trouble begins immediately once assertion is granted. If pure LLMs assert, they can testify; and if they can testify, we (human believers) can in principle acquire epistemic justification from their testimony, non-human testifiers though they are. Nothing about this is objectionable in general. A hearer who learns the boiling point of water from a machine has gained knowledge, and no epistemology of testimony worth the name should say otherwise. The trouble is local, and §4 is about where it is local to: testimonially based knowledge of other minds, in the one case where testimony is all there is.

4 The Argument from Other Minds

How do we know that another creature has a mind, not merely that it is intelligent, but that there is something it is like to be it? The traditional skeptical problem here trades on mere possibility: nothing rules out a world of zombies, therefore knowledge is threatened. Our argument trades on nothing of the kind, and it borrows none of the resources built to answer that problem.² Simonelli is not merely failing to rule out the zombie possibility. He gives us a positive, theoretically motivated argument that interlocutors occupying our communicative environment right now, or very shortly, satisfy the profile: full discursive competence, no sentience. A possibility we cannot exclude is one thing; a thesis we are asked to believe, on independent theoretical grounds, is another.³

²Avramides (2023) distinguishes a “thin” skeptical problem (that I am sometimes wrong about what a particular person thinks or feels) from a “thick” one pressing the Cartesian possibility that everyone but me is an automaton. It is the thick problem that is distinctively philosophical, and the standard replies to it (the argument from analogy, inference to the best explanation, criterial accounts, disjunctivist perceptual theories) are unaffected by anything Simonelli argues, or by anything we argue here.

³This forestalls the cheapest reply. Gomes (2019, §1) considers and sets aside cases of just this shape (“[I]sn’t it possible that they are zombies or robots, simulating desire and frustration but with nothing mental going on at all?”) classifying them under the *Problem of Error*, which “doesn’t turn on anything distinctive about our relationship to others’ minds.” Our reply is that inability to rule out a scenario and a theoretically motivated claim that the scenario is actual differ in epistemic kind, and only the first is what the Problem of Error sets aside. A second version of the reply comes from the epistemology of testimony, which already debates whether a “zombie testifier” (a source that “entirely lack[s] phenomenology”) can support testimonial knowledge, and tends to answer yes (Green, 2006, as cited in Green, n.d., §3.b.ii). But that debate asks whether epistemic *gain* is possible via such a source; ours asks whether the evidential relation between discursive performance and the performer’s own sentience survives such sources. An affirmative answer to the first supports a negative answer to the second. We lack direct access to Green’s 2006 dissertation and cite it only through the secondary literature that quotes it.

Consider, then, a purely text-mediated interlocutor, a lifelong pen pal, known through nothing but a lifetime of letters. Our confidence that such a person is phenomenally conscious rests centrally on the sophistication of their discursive performance, and on each leading account of how it could rest there, what SwS removes from a pure LLM is exactly what the account requires. If testimony is a basic source of knowledge of other minds, what lets a speaker's testimony *originate* such knowledge rather than merely relay it is her own non-testimonial access to her mental life, since "each subject has a non-testimonial way of knowing of her own engagements" (Gomes, 2015, p. 14), and a pure LLM has no such states to be acquainted with. If instead testimonial knowledge rests on the hearer's grasp of *expressive* behavior, what makes such behavior evidentially good is that it occurs "during an interval of time during which the person exemplifying the behaviour also instantiates the mental state being expressed" (Parrott, 2026, §1), and a pure LLM has no state for any interval to be an interval of. If what justifies the hearer is externalistically the speaker's own antecedent mental state (Munro, 2022), a pure LLM supplies none. We need not adjudicate among these accounts; on each of them, SwS removes precisely what the account needs. Two further observations reinforce the point independently of all three. Ordinary practice attributes phenomenal states at a granularity (a patient rating pain on a ten-point scale, or describing exactly what a hallucination looked like) that behavior alone radically underdetermines. And where no channel but text exists at all, discursive performance is not one evidential source among several; it is the only one there is. None of this touches the uncontroversial attributions we make to animals and preverbal infants on the strength of embodied behavior. Those rest on a different evidential base, and we return in defense of (OM5) to why the difference matters.

That evidential weight rests, in turn, on a background generalization: producing discursive performance of this sophistication is reliable evidence that the producer is sentient. The generalization never licensed the phenomenal attribution by itself. It licensed it in concert with a second fact, that the only producers of discourse at this level were human beings, of whom we independently take it that there is something it is like to be them. SwS leaves the second conjunct untouched and falsifies the first. One broken conjunct is enough to break the conjunction.

The Argument from Other Minds

(OM1) A pure LLM is a fully competent discursive performer and lacks phenomenal consciousness. [Simonelli's thesis, granted]

(OM2) If (OM1), then discursive performance is not *by itself* evidence that its producer has phenomenal consciousness.

(OM3) Prior to SwS the background generalization underwriting discursively mediated attribution was true, so text-mediated attributions of phenomenal consciousness could be knowledge.

(OM4) SwS falsifies that generalization while leaving the attributor's evidence untouched, so a true attribution formed under it is true by luck and therefore not *knowledge*, whether or not it remains justified.

(OM5) No non-discursive evidence available within Simonelli's own framework can restore it.

(OMC) Therefore, Simonelli must either restrict the SwS thesis or accept skepticism about *knowledge* of the phenomenal consciousness of text-mediated other minds. [from OM3–OM5]

In defense of (OM2). The severance is not an oversight. On Simonelli's own account, full discursive performance is constituted by two things and two things only: mastery of inferential role (§2), and the capacity to bear theoretical responsibility for what one says, answering challenges, clarifying, justifying, retracting under counter-evidence. He is explicit that this responsibility can be "siphoned" out of the fuller, practical responsibility a human speaker bears, leaving a notion of responsibility "in principle attributable to LLMs like ChatGPT" (Simonelli, 2026, p. 27). Neither component makes any reference to phenomenal consciousness, and the whole point of SwS is that neither could, on pain of collapsing back into the picture Simonelli rejects. There is, in his own theory, no constitutive link between discursive performance and phenomenal consciousness. This is not a gap he overlooked; it is the thesis. Note the qualification "by itself." Something can be evidence of something else without a constitutive link. Fingerprints are evidence of identity. What (OM2) denies is only the constitutively licensed inference; whatever contingent, statistical evidential force performance retains comes precisely from the background generalization in (OM3), and dismantling that generalization is the work of (OM4).

In defense of (OM3). The two preceding paragraphs defend this directly for the discursively mediated case: on any of the three leading accounts of testimonial and expressive knowledge, and independently of all of them in the two practices where discursive evidence carries irreplaceable weight.

In defense of (OM4). We state (OM4) as a claim about knowledge rather than justification deliberately, because it is both easier to defend and harder to escape. Justification is a matter of the attributor's evidential state; knowledge additionally requires that she not merely land on the truth by luck. After SwS the attributor's evidence does not change. She reads the same sophisticated, apparently sincere prose from her pen pal

that she always did. What changes is the environment that evidence sits in, now populated with indistinguishable producers of the same prose for the same kind of reasons, only without anything it is like to be them. A true belief formed in such an environment is true by luck, in the way familiar from cases where a subject's belief is correct but could too easily have been otherwise (Goldman, 1976, pp. 772–773; cf. Pritchard, 2005). This diagnosis has two advantages over a bare justification-loss claim. First, it does not require showing that the attributor's evidence has degraded: Simonelli can insist (and, given his own view of discursive performance, must insist) that the pen pal's letters remain exactly as good evidence of sapience as they ever were, and (OM4) survives that concession untouched, because it is a claim about the environment rather than about the evidence. Second, it blocks the reply that this is merely the old hard problem of other minds under a new name. The traditional problem concerns what conceptual or perceptual resources are in principle available; it is not sensitive to how common lookalikes happen to be in one's surroundings. What makes text-mediated attribution newly precarious is a fact about the composition of the communicative environment, and it is SwS, not the traditional problem, that is responsible for that fact.⁴

Each of these moves rests on the same single premise: that SwS is actual, or close enough to actual that non-sentient discursive performers genuinely populate our communicative environment, rather than a merely conceivable possibility. This is not a weakness we are hiding. It is the premise Simonelli himself supplies, and our whole strategy of immanent critique depends on taking him at his word about it (§1). If that premise is denied, the moves lose their force together, but denying it is denying SwS itself, which is not a position available to Simonelli.

In defense of (OM5). Non-discursive evidence of another's mind can come from only three places: from the interlocutor's *body*, from the *etiology* of what they produce, or from unmediated *acquaintance*, a fact about the other, a fact about the production, or a relation that bypasses evidence altogether. There is no fourth place, and within Simonelli's framework all three are closed.

Body. Unavailable by stipulation in the text-mediated case: the pen pal is known through nothing but letters, and a pure LLM is neither physically nor virtually embodied (§2). Nor would this exit succeed even where embodiment is available, to a designer deciding what to build, rather than to a correspondent deciding what to believe. Sen-

⁴A disjunctivist about expressive behavior (Parrott, 2026) might reply that a subject in the good disjunct (genuinely receiving an expression of a mental state) is not epistemically worse off for the existence of bad disjuncts elsewhere. But (OM4) grants the good disjunct: it grants that the pen pal is sentient and that her letters are genuinely, factively expressive. Its claim is that a true belief formed there still fails to be knowledge when the surrounding environment is saturated with subjectively identical bad disjuncts. Disjunctivism protects the good case against an argument that infers downward from indistinguishability; it was never designed to protect a true belief from bad epistemic luck in its surroundings.

tience in Simonelli's sense is, by his own admission, necessary but not sufficient for phenomenal consciousness (§2); confirming that a system instantiates a perception-action nexus therefore confirms at most the necessary condition, never the further, ethically loaded fact. For a biological creature that gap is closed by resources SwS leaves untouched, common physiology and evolutionary descent license an inference from function to phenomenology that owes nothing to performance of any kind, which is exactly why our argument never touched the attributions we make to animals and infants. An engineered system enjoys no such warrant: its functional profile is built to specification, and specification can secure sentience in Simonelli's sense without thereby securing what only a performance-to-consciousness inference could show, exactly the inference SwS puts in doubt. This is also where a proposal to look for deep *computational* markers of sentience belongs, markers of global broadcast, say, rather than markers of behavior. Such a proposal does not open a fourth place to look; it refines the first, since a system's computational organization is a fact about what that system is. And it inherits the limitation of the first: however far interpretability advances, what it delivers is a functional fact, and a functional fact secures the necessary condition alone. Simonelli's own language concedes as much: an architecture built around the perception-action nexus, he allows, "could in principle" yield genuinely sentient systems (Simonelli, 2026, p. 16), a hedge that is hard to explain if he already possessed a settled way to tell.

Etiology. This route is wider than it first appears. Simonelli can consistently maintain an asymmetry (sapience-attribution is constituted by performance, sentience-attribution is an empirical and evidential matter) so the charge that an appeal to causal history "conflicts with his own performance-based criterion" does not by itself close the exit. What closes it is what remains once the exit is taken. To appeal to causal history is to concede that the justification of sentience-attribution answers to facts outside discursive practice: to a system's production goals, its training regime, its computational structure. But the systems at issue are precisely those whose training optimizes them to produce text that reads as sincere, so the right causal history is no more verifiable than the wrong one. The exit relocates the verdict outside discourse without delivering one, and that no verdict is presently forthcoming from outside is conceded by the precautionary literature itself (Birch, 2024, pp. 320–321). In the text-mediated case, causal history is not evidence available in principle; it is evidence we do not have.

Acquaintance. Foreclosed by Simonelli's own foundational commitment that every part of the conceptual content of every expression is effable (§2). Positing an acquaintance-based residue that could do work discursive performance cannot is precisely the retreat to quasi-inference that GSI was built to avoid.

An anticipated reply. Simonelli might respond that he never claimed our confidence in a pen pal's consciousness rests on anything ineffable, and that he is under no obligation to say what does ground it, since he has been explicit throughout that his concern is conceptual understanding, not the metaphysics of consciousness, and phenomenal consciousness is exactly the kind of thing about which, on his own admission, one can do little more than "emit an inarticulate sound" (Simonelli, 2026) if pressed for its content. This will not do. The argument does not ask Simonelli to explain what grounds our confidence in other minds. It asks him to notice what his own positive theory entails about what *cannot* ground it. His theory is not silent on this question; it has a determinate answer. On that theory, however sophisticated a system's discursive performance becomes, that sophistication tracks nothing beyond inferential mastery and theoretical responsibility. Properties his own thesis defines independently of phenomenal consciousness. A generalization that held before SwS was on the table (that sophisticated discursive performance is reliable evidence of sentience) is a generalization his own theory now tells us not to trust, in precisely the case where nothing else is available to trust instead. Declining to say what *does* ground other-minds attribution is not neutrality once one's own theory has told us, in detail, what does not.

5 Two Costly Options

The first way out is to restrict SwS: carve out phenomenal vocabulary ("pain," "sees red," "feels tired") as an exception to the claim that inferential-role mastery is sufficient for concept possession. The natural form of the restriction is to say that, for this vocabulary specifically, inferential mastery is not enough; something like acquaintance, or a language-entry rule linking use to experience, is also required. But this is just the reintroduction of the quasi-inferential rules Simonelli spent his 2023 and 2026 papers arguing GSI does not need (§2). Once reintroduced for phenomenal vocabulary, there is no principled way to contain them: the same considerations that motivated their exclusion from color and pain concepts motivate their exclusion from any concept whose application conditions plausibly involve the world's showing up to a subject in a determinate way, and GSI's claim to full generality (the claim that made SwS an interesting thesis in the first place, rather than a modest remark about mathematics and logic) no longer holds. It is worth recording in advance that this exit does not carry the cost we will assign to the second. Once acquaintance or a language-entry rule is required for phenomenal vocabulary, mastery in the use of that vocabulary becomes, once again, a mark of phenomenal consciousness; for the language-capable systems the norm is actually about, a means of identification returns. The second exit keeps generality and recovers that mark nowhere. This is why the two costs are costs of different kinds.

The second way out is to accept systematic skepticism about knowledge of the phenomenal consciousness of other minds. This is not merely revisionary; it is in tension with a further commitment Simonelli states in his own paper’s final section. He argues at length that we should not aspire to bring sentient systems into existence, because sentient systems (creatures with well-being of which they are aware) pose “serious safety” risks and ethical burdens that sapient-but-non-sentient systems do not (Simonelli, 2026, pp. 29–30). He hopes that the emerging alien intelligence of advanced AI “stays” sapient without being sentient (p. 31). That hope, and the prescription behind it, presuppose that we can tell, at least well enough to act on, which systems are sentient and which are not. One might reply that the prescription targets only sentience in his defined, functional sense, which an inspection of architecture settles, so that no epistemic problem arises at all. What blocks the reply is the reasoning he gives for the prescription. He offers two grounds, and they are of different kinds. The safety ground (that such a system might determine that the continued existence of human beings is not in its own interest and act accordingly) holds for functional goal-pursuit alone, and we grant it. The ethical ground does not. The worry that we would incur obligations to these beings and, given the context in which they would be created as products for our use, would almost certainly fail to meet them, requires that they be moral patients; and what secures moral patiency is not the functional fact of instantiating a perception-action nexus but well-being of which the system is aware, a phenomenal fact. Simonelli’s own sentences follow this order: having characterized sentient creatures as ones to whom a world shows up as practically significant given their aims and capacities, he writes immediately that it seems clear we do not want systems that have a well-being of which they are aware (Simonelli, 2026, pp. 29–30). What the norm targets is that awareness, not the functional profile, and it is precisely the fact for which he has admitted only a necessary condition. What §4 shows, extended by the point made there in defense of (OM5), is that Simonelli’s own thesis removes the only marks of phenomenal consciousness available for exactly the systems the prescription is meant to bind. A pure, disembodied LLM is not where the norm does any work: on his own architectural criterion he already knows such a system lacks sentience in his sense, and by his own necessary-condition admission knows it is not phenomenally conscious either. There is nothing here to misidentify. The norm was never for that case. It is for the architectures he flags as genuine candidates, systems built around a perception-action nexus, the kind he attributes to LeCun’s proposal with visible hesitation, allowing only that it “could in principle” yield genuine sentience (p. 16). For exactly those systems, architecture confirms the necessary condition and nothing more, a limitation the precautionary literature concedes even while hoping to overcome it (Birch, 2024, p. 321); the further fact

the norm is actually about remains exactly as inaccessible as it was for the pen pal. *Do not build sentient systems* is not idle for lack of a target. It is idle because, for the one class of target that matters, Simonelli's own thesis leaves him no way to tell whether he has built one. Simonelli cannot have it both ways: a thesis that severs performance from phenomenal consciousness, applied to the very architectures he is warning us about, undercuts the practical prescription that thesis was meant to support.

6 Conclusion

We have not argued that Simonelli's hyper-inferentialism is false. We have argued that its truth is more expensive than its author has priced in. Grant Global Strong Inferentialism, grant that training on linguistic data alone can in principle secure mastery of any concept's inferential role, grant that a pure LLM can thereby be sapient without a shred of sentience. Grant, in short, everything Simonelli asks us to grant. What follows is not merely that pure LLMs mean something when they produce text. It is that they assert, and can be relied on testimonially, in exactly the sense that licenses our taking their word for things. And once discursive performance is severed from sentience by thesis rather than by accident, the one channel through which we have always come to know that a purely text-mediated interlocutor is not a zombie (a lifelong pen pal, known through nothing but letters) stops being able to do that work. Not because some new skeptical scenario has been dreamed up: the traditional problem of other minds, and the traditional Problem of Error that a defender of Simonelli might try to assimilate our argument to, are both untouched by anything we have said (§4). What defeats the pen-pal attribution is not a possibility we cannot rule out but an environmental fact Simonelli's own thesis puts in place, and the resulting failure is properly a failure of knowledge, not merely of justification, which is why the usual replies available to a defender of testimony's basicness or of disjunctivism about expression do not restore it (§4).

Neither exit is free. Restricting SwS to exempt phenomenal vocabulary reintroduces the quasi-inferential, language-entry rules that GSI was built without, and with them goes the generality that made the thesis worth defending in the first place. GSI becomes, at best, a claim about mathematics, logic, and whatever else can be had without acquaintance, which is a much smaller and less interesting thesis than the one Simonelli set out to defend. Accepting skepticism about knowledge of other minds' phenomenal consciousness is not merely intuitively costly; it is in tension with the one practical commitment Simonelli explicitly stakes out, that we should not aspire to build sentient systems. A norm against building things we cannot identify is not a norm we can follow, and Simonelli's own thesis is what removes our ability to identify them, not for

the pure LLMs he has already classified by architecture, but for the embodied architectures he himself flags as genuine candidates, which is where the norm was supposed to bind. We suspect Simonelli will not want to give up the generality of SwS, and we are confident no one, including Simonelli, wants to give up ordinary confidence that the people we talk to are not zombies. Both, on his own theory, cannot be kept.

This work may bring out one further thing. Many earlier approaches in the philosophy of language, which did not separate the semantic thesis from the pragmatic one, were implicitly relying on the contingent fact that genuine language agents were humans like us (cf. Cappelen & Dever, 2021). What the philosophy of language for AI is revealing is twofold: that there can perfectly well be genuine language agents who are not humans like us, and that a semantic framework for them needs to be finer-grained. The distinction between meaning-expression and assertion that we have put to work in this paper is one way of making the framework finer-grained.

Hinton and LeCun disagree about whether current systems understand what they are saying. Simonelli's paper is best read as siding with Hinton and then some: not only do they understand, they may, in principle, understand everything there is to understand, all the way down to concepts of pain and color and grief, without any accompanying awareness whatsoever. We have argued that this is not a free-standing theoretical curiosity. It reaches back into the epistemic practices by which we know one another, and it reaches back precisely at the point where those practices have nothing else to stand on. The move that lets us gain knowledge from machines costs us knowledge of persons.

References

- Avramides, A. (2023). Other Minds. In E. N. Zalta & U. Nodelman (Eds.), *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University.
- Birch, J. (2024). *The Edge of Sentience: Risk and Precaution in Humans, Other Animals, and AI*. Oxford: Oxford University Press. doi: 10.1093/9780191966729.001.0001
- Borg, E. (2026). LLMs, Turing Tests and Chinese Rooms: The Prospects for Meaning in Large Language Models. *Inquiry*, 69(6), 2807–2837. doi: 10.1080/0020174X.2024.2446241
- Brandom, R. B. (1994). *Making It Explicit: Reasoning, Representing, and Discursive Commitment*. Cambridge, MA: Harvard University Press.
- Butlin, P., & Viebahn, E. (2025). AI Assertion. *Ergo*, 12, 968–988. doi: 10.3998/ergo.7960
- Cappelen, H., & Dever, J. (2021). *Making AI Intelligible: Philosophical Foundations*. Oxford: Oxford University Press. doi: 10.1093/oso/9780192894724.001.0001

- Gibson, J. J. (1979). The Theory of Affordances. In *The Ecological Approach to Visual Perception* (chap. 8). Boston: Houghton Mifflin.
- Goldman, A. I. (1976). Discrimination and Perceptual Knowledge. *The Journal of Philosophy*, 73(20), 771–791. doi: 10.2307/2025679
- Gomes, A. (2015). Testimony and Other Minds. *Erkenntnis*, 80(1), 173–183. doi: 10.1007/s10670-014-9619-8
- Gomes, A. (2019). Perception, Evidence, and Our Expressive Knowledge of Others’ Minds. In A. Avramides & M. Parrott (Eds.), *Knowing Other Minds* (pp. 148–172). Oxford: Oxford University Press. doi: 10.1093/oso/9780198794400.003.0008
- Green, C. R. (n.d.). Epistemology of Testimony. In *Internet Encyclopedia of Philosophy*.
- Hattiangadi, A., & Schoubye, A. J. (2025). The Outputs of Large Language Models Are Meaningless. In H. Cappelen & R. Sterken (Eds.), *Communicating with AI: Philosophical Perspectives*. Oxford: Oxford University Press.
- Mallory, F. (2023). Fictionalism about Chatbots. *Ergo*, 10. doi: 10.3998/ergo.4668
- Munro, D. (2022). Mental Imagery and the Epistemology of Testimony. *Canadian Journal of Philosophy*, 52(4), 428–449. doi: 10.1017/can.2022.42
- Parrott, M. (2026). Expression, Testimony, and Other Minds. *Philosophical Studies*, 183(1), 287–305. doi: 10.1007/s11098-025-02443-4
- Pritchard, D. (2005). *Epistemic Luck*. Oxford: Oxford University Press. doi: 10.1093/019928038X.001.0001
- Sambrotta, M. (2025). LLMs and the Logical Space of Reasons. *Minds and Machines*, 35(4), 46. doi: 10.1007/s11023-025-09751-y
- Simonelli, R. (2023). How to Be a Hyper-Inferentialist. *Synthese*, 202(5). doi: 10.1007/s11229-023-04382-1
- Simonelli, R. (2026). Sapience without Sentience: An Inferentialist Approach to LLMs. *Asian Journal of Philosophy*, 5(1). doi: 10.1007/s44204-026-00400-4
- Tiehen, J. (in press). Lack a Theory of Mind and So Can’t Perform Speech Acts: A Causal Argument. *Philosophy of AI*.